Project Title

Clothing Item Classification Using Deep Learning on Fashion MNIST

Project Overview

This project aims to develop a deep learning model using Fashion MNIST to classify clothing items into 10 categories (T-shirt, Trouser, Pullover, Dress, Coat, Sandal, Shirt, Sneaker, Bag, Ankle Boot). The model will be trained using a Convolutional Neural Network (CNN) implemented in Python and PyTorch.

Why is this project good?

- Fashion item classification is a fundamental problem in computer vision, applicable in e-commerce, retail, and automated inventory systems.
- Deep learning, especially CNNs, has been highly effective for image classification tasks.
- This project will enhance understanding of CNNs, feature extraction, and optimization techniques.
- The Fashion MNIST dataset is well-suited for testing and benchmarking classification models.

Key Objectives:

- Understand and preprocess the Fashion MNIST dataset (grayscale 28x28 pixel images).
- Build and train a CNN model to classify clothing items.
- Optimize model performance using hyperparameter tuning, dropout, and batch normalization.
- Compare CNN performance with fully connected neural networks (MLP).
- Visualize model performance using accuracy curves and a confusion matrix.

Project Methodology:

1. Data Preprocessing:

- Load the Fashion MNIST dataset.
- Normalize pixel values and reshape images.
- Perform data augmentation (rotation, zoom, shift) to improve generalization.

2. Model Implementation:

- Implement a CNN with multiple convolutional layers using PyTorch.
- Compare it with a *basic fully connected network (MLP).
- Train both models and evaluate their accuracy.

3. Performance Evaluation:

- Use accuracy, precision, recall, F1-score as evaluation metrics.
- Visualize training/validation loss and accuracy curves.
- Plot a confusion matrix to analyze misclassified images.

4. Optimization Techniques:

- Apply Dropout and Batch Normalization to improve generalization.
- Tune hyperparameters (learning rate, batch size, optimizer).
- Use data augmentation to enhance model performance.

5. Results & Interpretation:

- Compare CNN with MLP and discuss why CNN performs better.
- Highlight challenges in fashion item classification.

What data will be used?

The Fashion MNIST dataset will be used, which consists of 70,000 images (60,000 for training and 10,000 for testing). Each image is a 28x28 grayscale image belonging to one of 10 clothing categories.

How will the system's performance be evaluated?

- The system's performance will be measured using accuracy, precision, recall, and F1-score.
- A confusion matrix will be generated to analyze misclassifications.
- Training and validation loss and accuracy curves will be plotted to visualize the learning process.
- The performance of the CNN model will be compared to an MLP model to demonstrate the advantages of deep feature extraction.

Expected Outcome:

- A trained deep learning model achieving high classification accuracy on Fashion MNIST.
- Performance comparison between CNN and MLP architectures.
- Insights into model optimization and real-world applications (e.g., fashion retail automation).